Which fusion methods have actually been used?

Information / method	What operation is this?	Current implementation
Word + speaker intervals	Temporal overlap association	Implemented; audio-derived information
Temporal + semantic evidence	Weighted score-level fusion	Implemented; optional timing; 0.5 / 0.5
Word + character + context TF-IDF	Feature concatenation	Implemented inside text classifier
Forest + logistic probabilities	Classifier-output ensemble	Implemented inside text; 0.75 / 0.25
MiniLM versus TF-IDF	Alternative representations	Selectable; not fused together
Participation + contribution features	Research feature combination	Implemented in GAP ranking; no slide input
Early / intermediate audio–document network	Joint learned multimodal fusion	Not deployed; separate toy prototypes exist
Cross-attention / tensor / learned graph fusion	Learned interaction mechanisms	Not implemented
Learned reliability gating	Input-dependent weighting	Not implemented; fixed missing-cue rule only

The temporal–semantic hybrid combines cues at the score level. It is not an architecture combining early and late neural fusion. Within-text concatenation/ensembling does not create extra input modalities.