The actual alignment fusion rule and missing-timing fallback

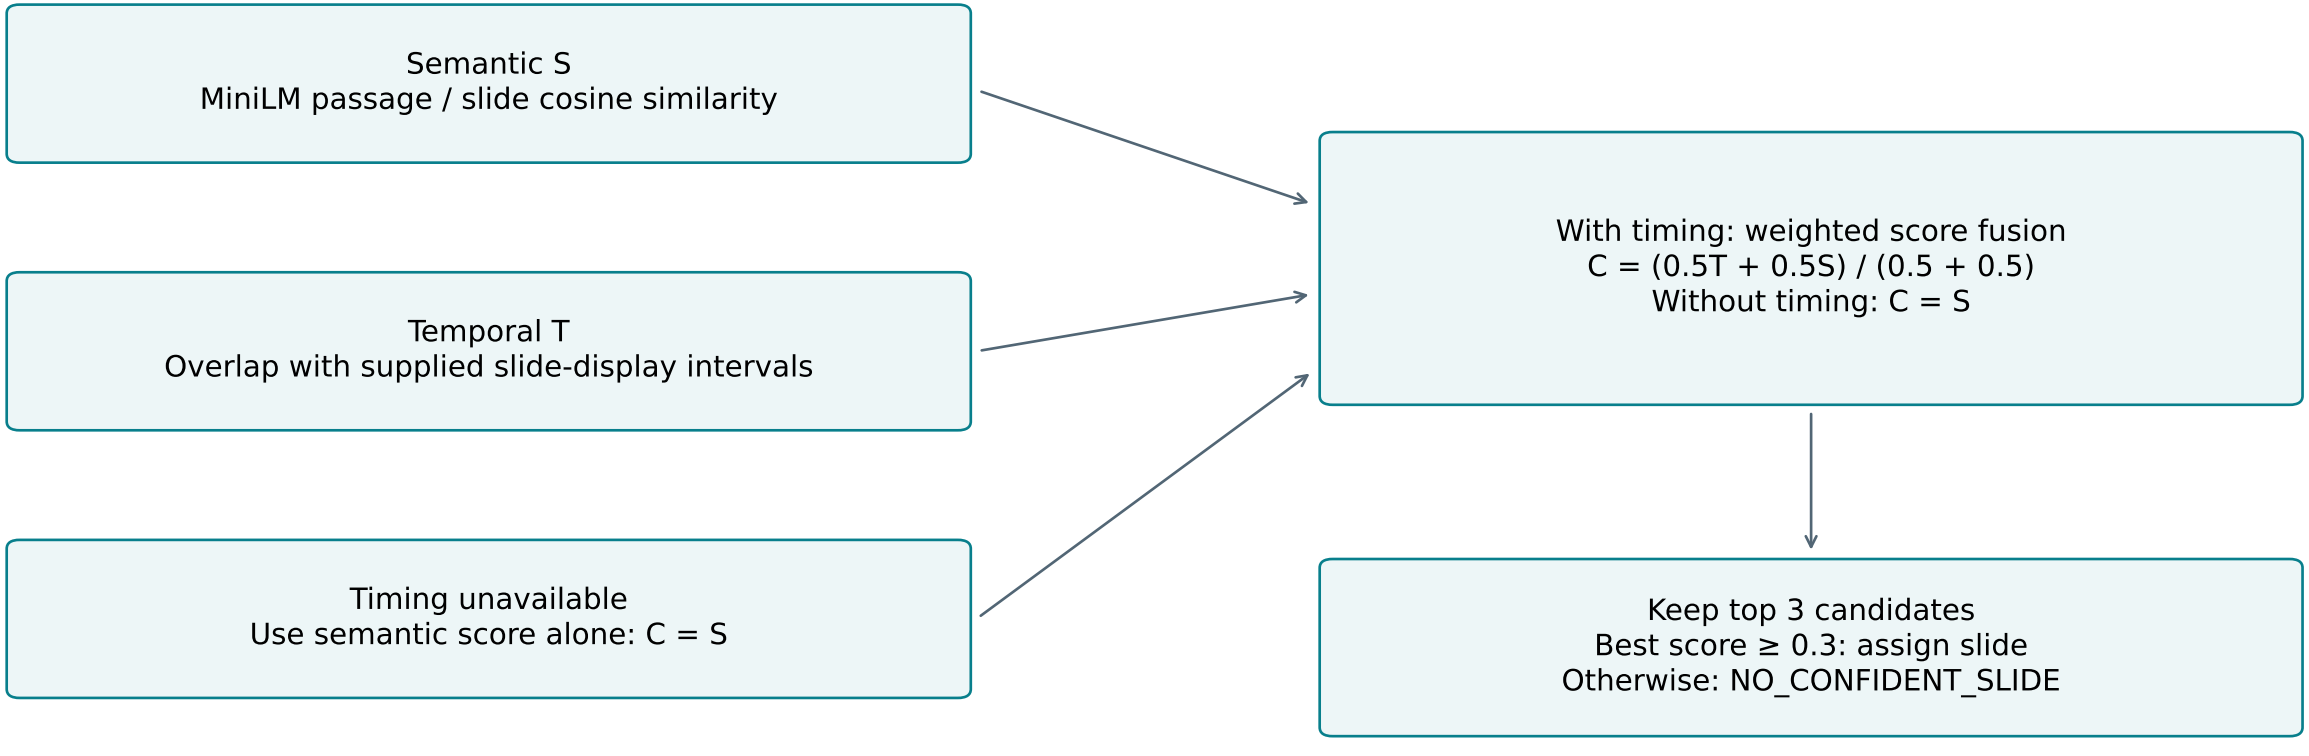

Configuration/implementation diagram, not measured accuracy. Pure temporal-only and semantic-only settings serve as single-cue baselines. Fixed weighted cue fusion is not learned cross-attention or a neural early-plus-late hybrid architecture.